

TRIAI x Google DeepMind's AI Research foundations course

Challenge 4: Mini-Engagement Plan

Project title: Early Detection of Toxic Language through Ubuntu-Aligned Latent Probing of LLMs

Participants of the project

Birhanu Simachew

Semere Herruy

Gebrhans Weldegebriel

1. Stakeholder Groups Selected

1.1. Online Community Moderators & Civil Society Digital Safety Advocates

1.2. Youth Social Media Users

2. What Matters Most

Community Moderators and advocates prioritize harm reduction while safeguarding freedom of expression. They want AI systems that reduce workload, integrate into moderation workflows, and avoid over-flagging marginalized voices. Transparency in dataset sourcing and accountability in deployment are critical. They value tools that empower communities to self-govern digital spaces, aligned with Ubuntu's ethic of collective responsibility. Their goal is to ensure AI moderation strengthens trust and safety without silencing legitimate cultural or political discourse.

Youth social media users care about safe participation online without harassment or exclusion. They want AI systems that understand slang, code-switching, and local languages, ensuring inclusivity across rural and urban contexts. Their priority is fairness: toxicity detection should not silence playful or culturally specific speech. They value dignity, creativity, and empowerment, seeking tools that make online spaces safer while respecting linguistic diversity. They want assurance that AI reflects their lived realities and supports Ubuntu's principle of community belonging.

3. Potential Harms & Benefits

Stakeholder	Potential Harms	Potential Benefits
Moderators & Advocates	Over-flagging legitimate speech, exclusion of minority dialects, lack of transparency	Reduced workload, safer communities, stronger trust
Youth Users	Misinterpretation of slang, silencing of cultural expression, bias against rural dialects	Safer participation, empowerment, inclusive representation

4. Engagement Methods (Action Catalogue)

Moderators & Advocates → User Committees where moderators and advocates review flagged outputs, propose dataset refinements, and co-design harm-reduction strategies. Ensures accountability and transparency.

User Committees with Moderators & Advocates institutionalizes community oversight. Moderators and advocates are already embedded in digital safety work, so their participation ensures probe design reflects real harms and practical needs. Committees provide structured feedback loops, reviewing flagged examples and advising on dataset gaps. This builds accountability, prevents misuse, and aligns with Ubuntu by fostering collective responsibility. It also ensures transparency in how toxicity detection decisions are made, strengthening trust between AI developers and communities.

Youth Users → Interactive workshops where youth share experiences of toxicity, demonstrate slang/code-switching, and evaluate model outputs. Captures diverse linguistic realities and empowers youth voices.

Interactive workshops allow young people to articulate how toxicity manifests in their daily online interactions. They can demonstrate slang, code-switching, and multilingual expressions that AI systems often misinterpret. By reviewing outputs together, youth participants highlight cultural nuances and biases. This method empowers them as co-designers, ensuring inclusivity and fairness. It also surfaces risks of silencing playful or creative speech, guiding dataset curation to avoid bias. Ubuntu ethics are embedded here by valuing youth voices and treating them as equal partners in shaping safer digital spaces.

Together, these methods ensure that toxicity detection is not imposed top-down, but rather co-created with those most affected. They mitigate risks of exclusion and misrepresentation while maximizing benefits like empowerment, equity, and cultural preservation. This approach operationalizes Ubuntu ethics: *“I am because we are.”*

5. Influence on Dataset & Model Design

- User Committees: Identify false positives and overlooked harms, ensuring balanced representation across dialects and contexts.

- Youth Workshops: Enrich datasets with slang, code-switching, and multilingual examples, improving cultural sensitivity.

The result is a toxicity detection system that is Ubuntu-aligned: inclusive, transparent, and responsive to community needs, ensuring AI serves as a tool for empowerment rather than exclusion.